ProPublica's COMPAS Data Revisited

Matias Barenstein

July 08, 2019

Abstract

I examine the COMPAS recidivism risk score and criminal history data collected by ProPublica in 2016 that fueled intense debate and research in the nascent field of 'algorithmic fairness'. ProPublica's COMPAS data is used in an increasing number of studies to test various definitions of algorithmic fairness. This paper takes a closer look at the actual datasets put together by ProPublica. In particular, the sub-datasets built to study the likelihood of recidivism within two years of a defendant's original COMPAS survey screening date. I take a new yet simple approach to visualize these data, by analyzing the distribution of defendants across COMPAS screening dates. I find that ProPublica made an important data processing error when it created these datasets, failing to implement a two-year sample cutoff rule for recidivists in such datasets (whereas it implemented a two-year sample cutoff rule for non-recidivists). When I implement a simple two-year COMPAS screen date cutoff rule for recidivists, I estimate that in the two-year general recidivism dataset ProPublica kept over 40% more recidivists than it should have. This fundamental problem in dataset construction affects some statistics more than others. It obviously has a substantial impact on the recidivism rate; artificially inflating it. For the two-year general recidivism dataset created by ProPublica, the two-year recidivism rate is 45.1%, whereas, with the simple COMPAS screen date cutoff correction I implement, it is 36.2%. Thus, the two-year recidivism rate in ProPublica's dataset is inflated by over 24%. This also affects the positive and negative predictive values. On the other hand, this data processing error has little impact on some of the other key statistical measures, which are less susceptible to changes in the relative share of recidivists, such as the false positive and false negative rates, and the overall accuracy.

Keywords: Fair Machine Learning, Algorithmic Fairness, Recidivism, Risk, Bias, COMPAS, ProPublica[1]

---

1   Introduction

Due to the rise in data collection and its use, as well as the accompanying development of more predictive and complex machine learning models, and the still nascent field of artificial intelligence, the past several years has seen a marked spike in interest and research regarding what is now often referred to as "algorithmic fairness" or "fair machine learning." (See Corbett-Davies and Goel 2018; Cowgill and Tucker 2019; Kleinberg et al. 2018)[2]

One notable event in the chronological development of this field, which helped propel the interest and research into algorithmic fairness, was the ground-breaking investigative journalism work of ProPublica on the COMPAS recidivism risk score, which is sometimes used to aide various decisions in the judicial system.[3] In 2016, a team of journalists from ProPublica constructed a dataset of defendants from Broward County, FL, who had been arrested in 2013 or 2014 and assessed with the COMPAS risk screening system. ProPublica then collected data on future arrests for these defendants through the end of March 2016, in order to study how the COMPAS score predicted recidivism (Angwin et al. 2016).

Based on its analysis, focusing on one set of predictive metrics, ProPublica concluded that the COMPAS risk score was biased against African-Americans. The company that developed the COMPAS risk scoring system, Northpointe Inc., focusing on a different set of predictive metrics, defended the risk scores as unbiased (Dieterich et al. 2016). This sparked intense debate and research on the various possible definitions of algorithmic fairness. Some of this work has been primarily theoretical. For example, the work showing the impossibility of simultaneously attaining some of the more popular fairness goals (Chouldechova 2016; Kleinberg et al. 2018).

ProPublica's investigation was truly ground-breaking since it used public records requests to obtain COMPAS scores and dates and personal information on a group of defendants, as well as prison and jail information for them, and was able to match and merge these disparate data sources. Moreover, perhaps for full transparency and following best practices for reproducibility, ProPublica made the dataset it collected available to the public. As a result, the ProPublica COMPAS data has become one of the key bench-marking datasets with which a growing number of researchers have tested novel algorithmic fairness definitions and procedures. Indeed, it has become perhaps the most widely used dataset in the field of algorithmic fairness (e.g. Corbett-Davies and Goel 2018; Bilal Zafar et al. 2016, 2017; Chouldechova 2016; Corbett-Davies et al. 2017; Cowgill 2018; Rudin et al. 2018).

While ProPublica's COMPAS data is used in an increasing number of studies, researchers have taken the datasets created by ProPublica as they are and do not appear to have examined them closely for data processing issues.[4] Instead of testing a novel fairness definition or procedure, I take a closer look at the actual datasets put together by ProPublica. In particular, the sub-datasets built to study the likelihood of recidivism within two years of the original offense and COMPAS screening date. I take a new yet simple approach to visualize these data, by analyzing the distribution of defendants across COMPAS screening dates. Doing so, I find that ProPublica made an important data processing error creating some of the key datasets most often used by other researchers.

ProPublica failed to implement a two-year sample cutoff rule for recidivists in the two-year recidivism datasets (whereas it implemented a two-year sample cutoff rule for non-recidivists in the same datasets). As a result, ProPublica incorrectly kept a disproportionate share of recidivists in such datasets. This data processing mistake leads to biased two-year recidivism datasets, with artificially high recidivism rates. To my knowledge, this is the first paper to highlight this key data processing mistake.

---

[2]See also a seminal paper in this literature by Barocas and Selbst (2016).

[3]A private company, Northpointe Inc. (now Equivant), developed the COMPAS recidivism risk scores. COMPAS is short for Correctional Offender Management Profiling for Alternative Sanctions.

[4]Except for Rudin et al. (2018), who reconstruct datasets from the original ProPublica Python database "partly to ensure the quality of the features and partly to create new features." (p.32) While in so doing they may have avoided making the same data processing mistake that ProPublica makes, they do not generally highlight the dataset differences between their data and ProPublica's and do not identify ProPublica's data processing mistake. Their focus is altogether different, as they attempt to reverse engineer the COMPAS recidivism risk scores, to understand how Northpointe builds those scores.

To construct these two-year recidivism sub-datasets, ProPublica presumably wanted to keep people observed for at least two years at the end of the time window for which ProPublica collected criminal history data, on April 1, 2016. Therefore, we should not have expected to see anybody in the two-year datasets with COMPAS screening (or arrest) dates after April 1, 2014 (i.e. less than two years prior to ProPublica's data collection). However, as we will see further below, there are many people in ProPublica's two-year recidivism datasets who do indeed have COMPAS screening dates after this potential cutoff, all the way through December 31, 2014, which is the end date of the original database.

Taking a closer look at these datasets, I find that ProPublica correctly dropped non-recidivists with COMPAS screening dates post 4/1/2014. However, it kept people with COMPAS screening dates after 4/1/2014 if they recidivated. ProPublica's data processing logic that created the two-year recidivism datasets is as follows: keep a person if they are observed for at least two years OR keep a person if they recidivate within two years. Unfortunately, this results in a biased sample for the two-year recidivism datasets.[5]

As I show in this paper, the bias in the two-year datasets is clear, there are a disproportionate number of recidivists. When I implement a simple two-year COMPAS screen date cutoff rule for recidivists, I estimate that in the two-year general recidivism dataset ProPublica kept over 40% more recidivists than it should have. This fundamental problem in dataset construction affects some statistics more than others. It obviously has a substantial impact on the recidivism rate. It artificially inflates the recidivism rate. For the two-year general recidivism dataset created by ProPublica, the two-year recidivism rate is 45.1%, whereas with the simple COMPAS screen date correction I implement, it is 36.2%. Therefore, the two-year recidivism rate in ProPublica's two-year dataset is inflated by over 24%. This also affects the positive predictive value (PPV) or precision, and the negative predictive value (NPV). On the other hand, it has relatively little impact on several other key statistics that are less susceptible to changes in the relative share of recidivists versus non-recidivists, such as accuracy, the false positive rate (FPR), and the false negative rate (FNR).[6]

In the remainder of this paper, Section 2 briefly discusses the data collected by ProPublica. Section 3 examines in detail the data processing mistake just highlighted. Section 4 shows how this mistake artificially inflates the two-year general recidivism rate. In Section 5, I replicate ProPublica's confusion matrix analysis with the original two-year data, and then show the analogous results with the corrected COMPAS screen date cutoff, finding that in addition to the recidivism rate (or prevalence), the PPV, NPV, and detection rate, are also substantially impacted, but the FPR, FNR, and accuracy are not. Section 6 concludes. The Appendix discusses various key data and analysis assumptions, as well as some extensions.

2 Data

ProPublica obtained a dataset of pretrial defendants and probationers from Broward County, FL, who had been assessed with the COMPAS screening system between January 1, 2013, and December 31, 2014. COMPAS recidivism risk scores are based on a defendant's answers to the COMPAS screening survey. The survey is completed by pre-trial services in cooperation with the defendant after his or her arrest.[7] The COMPAS survey, at least in the ProPublica data, is typically administered the same day or the day after a person is jailed.[8]

For the more than 11 thousand *pretrial* defendants in this dataset,[9] ProPublica then collected data on future

---

[5]It is not clear whether ProPublica intended to actually process the data this way, in which case it is a conceptual mistake, or whether it did not intend to use this faulty logic, in which case it is a data processing mistake. In either case, it leads to the same biased sample two-year datasets.

[6]Or one minus these rates, i.e. specificity and sensitivity. Note that the FPR, and hence, also specificity, are by definition independent of the actual number of positives (or recidivists) in the data. They are calculated based only on actual negatives (i.e. only people who do *not* recidivate). So these statistics remain unchanged when I implement a two-year sample cutoff on recidivists.

[7]The COMPAS survey contains over 130 questions; part of the survey is based on administrative data (Cowgill 2018).

[8]In ProPublica's COMPAS data 69% of defendants appear to be administered the COMPAS survey the same day or one day after they are jailed.

[9]The set of more than 11 thousand pretrial defendants is what I call the *full* dataset. It has 11,757 people. (I also call the slightly smaller set of 10,331 people the *full* dataset; this second variant of the full dataset is reduced due to the full dataset

arrests through the end of March 2016, in order to study how the COMPAS score predicts recidivism for these defendants.[10]

ProPublica collected the data for its study and created a (Python) database. From that database, it constructed various sub-datasets that merged and calculated various important features. For example, an indicator for a re-arrest for a new crime within two years of the original one, and the period of time between arrests. ProPublica then exported these sub-datasets into .csv files.[11] These files are the ones most often used by other researchers (see references in the introduction).

I primarily use two of the .csv files that ProPublica created. These were named by ProPublica "compas-scores.csv" and "compas-scores-two-years.csv". The first file contains the full dataset of *pretrial* defendants that ProPublica obtained from the Broward County Sheriff's Office. This file contains 11,757 people.[12] This total is trimmed down by ProPublica to 10,331 people (I discuss the trimming done by ProPublica in the Appendix).

The second file I use is a file that ProPublica created for the purpose of studying two-year general recidivism. The term *general* recidivism is used to distinguish it from the smaller subset of *violent* recidivism. General recidivism includes both violent and non-violent offenses. I focus on the two-year general recidivism dataset in this paper, but the two-year violent recidivism data created by ProPublica has the same data processing issue. These files contain, in theory, a subset of people who are observed for at least two years, and it tags people who recidivated within two years as having a *two_year_recid* indicator flag turned on.[13] The two-year general recidivism file contains 7,214 people.

3   ProPublica's Data Processing Mistake

I start by looking at the full dataset of 10,331 people. I graph the number of cases or arrests by COMPAS screening date (and draw a vertical red line on April 1, 2014, which is the point in time after which we should not see any people entering the more reduced two-year recidivism sub-dataset(s) that I graph further below, as just explained). For this, and the subsequent histograms of COMPAS screening dates, I use 7-day (i.e. week-long) data bins.

---

trimming done by ProPublica, which I describe in the Appendix)

[10]ProPublica obtained criminal history information (both before and after the COMPAS screen date) for this sample of COMPAS pretrial defendants from public criminal records on the Broward County Clerk's Office website through April 1, 2016. It also obtained jail records from the Broward County Sheriff's Office from January 2013 to April 2016 and downloaded public incarceration records from the Florida Department of Corrections website.

[11]The location of the ProPublica data on the Web is at https://github.com/propublica/compas-analysis.

[12]That file does not contain some key information in other datasets, such as any prison time served for the original crime if convicted, which is necessary to calculate whether the person was free (outside prison) for at least two years, and thus, it does not have a flag for two-year recidivism.

[13]In addition to the distinction between general and violent recidivism, I sometimes draw a distinction between 'overall' or 'any' recidivism versus two-year recidivism. This distinction is not based on the type of offense committed, but rather on the timing of the offense, as I explain in the next section.
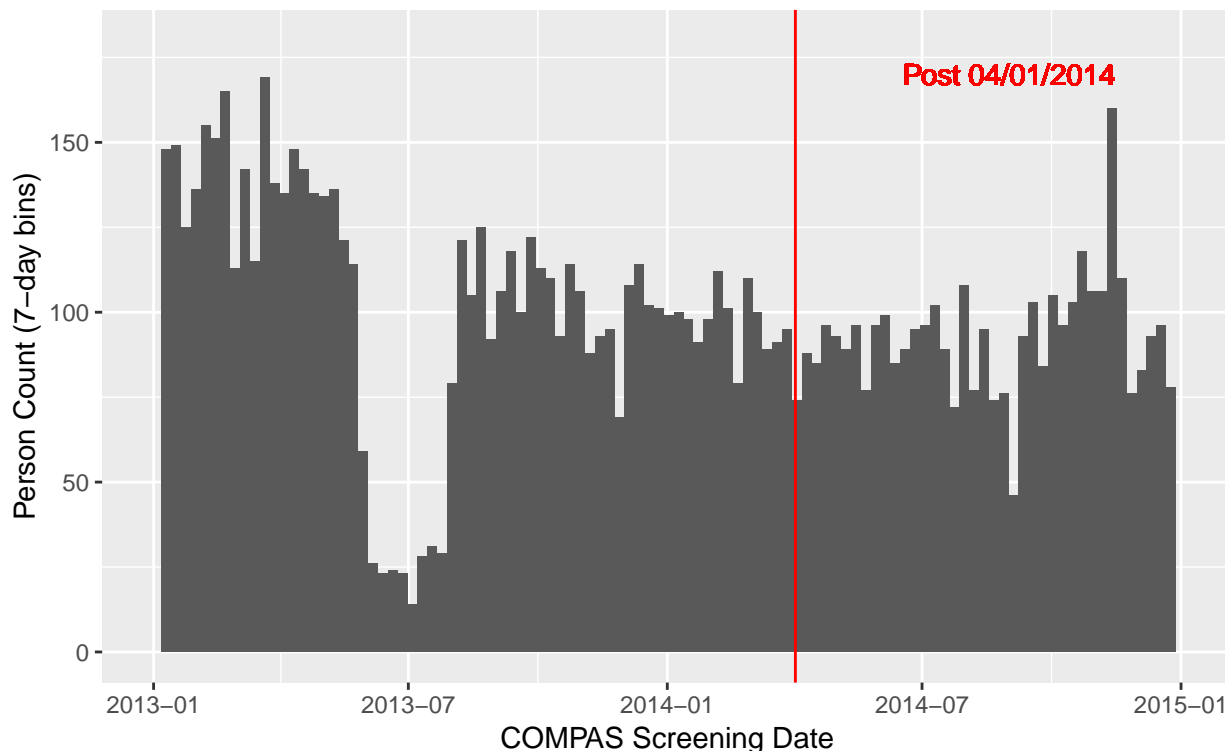
Figure 1: Persons by COMPAS Screen Date (7-day bins) - ProPublica Full Dataset

Other than the very noticeable drop in COMPAS screen dates in mid-2013, this graph appears reasonable.[14] The dates where the mid-2013 drop occurs are in June and July 2013. It is not clear why there is such a drop in COMPAS cases during these two months. I do not address this issue in this paper. (To the extent this is a problem, it appears to be a problem with the original dataset that ProPublica received from Broward County since it is also evident in ProPublica's "compas-scores-raw.csv" dataset. So it does not appear to be a data processing issue by ProPublica)[15]

To construct the two-year recidivism dataset(s), ProPublica presumably wanted to keep people observed for at least two years at the end of the time window for which it collected criminal history data, on April 1, 2016. As mentioned in the introduction, we should not have expected, therefore, to see *anybody* in the two-year datasets with COMPAS screening (or arrest) dates after 4/1/2014 (i.e. less than two years prior to ProPublica's data collection). However, as I show in the next Figure, there are many people in the two-year dataset who do indeed have a COMPAS screening (or arrest) date after this potential cutoff, all the way through December 31, 2014 (which is the end date of the full database).

---

[14]Also noticeable is the generally higher number of COMPAS screen dates in the first half of 2013.

[15]I checked whether the relatively few people with COMPAS screen dates during these two months looked different than the rest of the data, along various dimensions, but they generally did not, except they did have somewhat longer times between the arrest date and COMPAS screen date; see Appendix.

Persons by COMPAS Screen Date
### ProPublica Two–Year Dataset [N = 7214]

Figure 2: Persons by COMPAS Screen Date (7-day bins) - ProPublica Two-Year Dataset

The potential two-year COMPAS screen cutoff mark is indicated by the red vertical line on April 1, 2014. We clearly see that while the number of people in the two-year (general recidivism) dataset drops substantially after 4/1/2014, there are still non-trivial numbers of people after that date. This is because, as mentioned above, there was an error in ProPublica's data processing used to create the two-year recidivism datasets.

To create the two-year dataset, ProPublica used the following logic. You either had your COMPAS screen date at least two years prior to ProPublica's data collection time. So two years prior to the end of March 2016 (net of any jail and prison time). Or you could be in the data for less than two years if you recidivated. Unfortunately, for the latter, ProPublica did not use the cutoff of 4/1/2014 for the COMPAS screen date. So this creates an unbalanced dataset with too many recidivists. This is shown more clearly in the Figures below.

To see the data processing mistake more clearly, I now take a look at these COMPAS screen dates separately for recidivists and non-recidivists. I show first ProPublica's full dataset, and then ProPublica's two-year general recidivism dataset. For ease of comparison, I do this for the overall or any recidivism variable (i.e. the "is_recid" variable in ProPublica's dataset; instead of the "two_year_recid" variable).[16]

---

[16]I do not use two-year recidivism here since the full dataset does not have a two-year recidivism flag. As we see in the Recidivism Rates Section below, there are 220 people who have the overall "is_recid" flag turned on, but not the "two_year_recid" flag. These are people who recidivated but did so after more than two years after the original COMPAS screen date, but before the end date of ProPublica criminal history data window, at the end of March 2016. These 220 people represent only a 0.06 share of the 3,471 people who recidivate in total.

Persons by COMPAS Screen Date by Recidivism Status
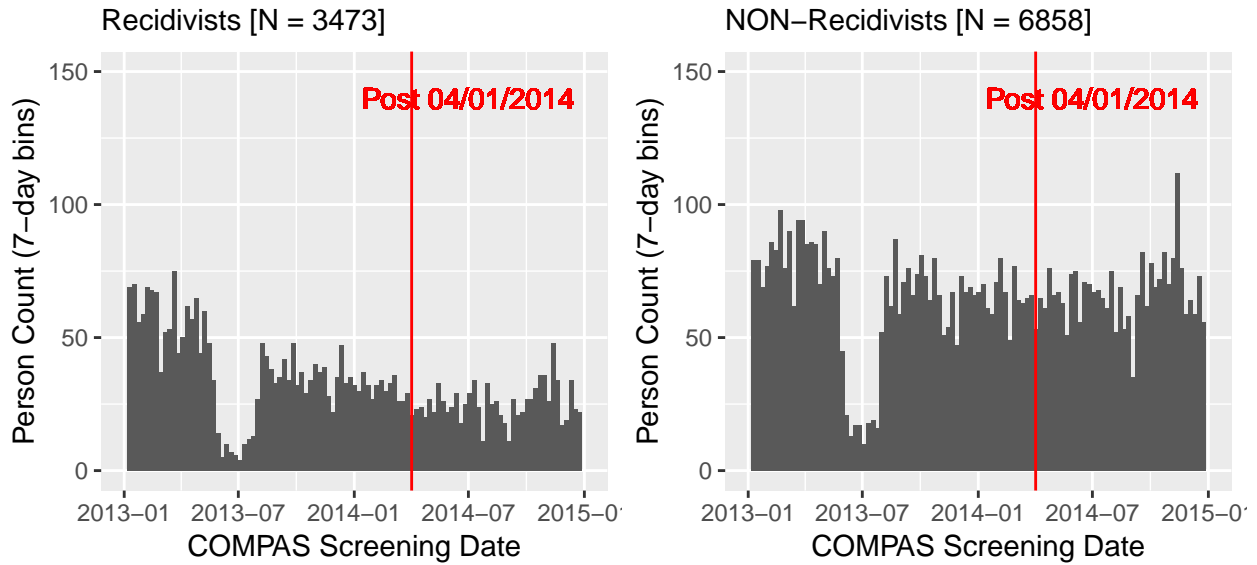**[ProPublica Full Dataset]**

Figure 3: Persons by COMPAS Screen Date (7-day bins) by Recividism Status - ProPublica Full Dataset

Persons by COMPAS Screen Date by Recidivism Status
[ProPublica Two−Year Dataset]

Figure 4: Persons by COMPAS Screen Date (7-day bins) by Recidivism Status - ProPublica Two-Year Dataset

The graphs on the top panel are as expected. There is no real difference in the pattern between recidivists and non-recidivists (other than the overall level). The bottom panel graphs for the two-year dataset, however, do show a stark difference between recidivists and non-recidivists. We see that while ProPublica correctly dropped all non-recidivists with COMPAS screening dates post 4/1/2014 in the two-year dataset, it kept

people with COMPAS screening dates after 4/1/2014 in that data if they recidivated. Indeed, in the Tables below we see that the two-year recidivism dataset has almost exactly the same number of people who recidivate at any point in time, as the full data does, 3,471 vs. 3,473.[17]

Table 1: Any Recidivism - ProPublica Full Dataset vs. ProPublica Two-Year Dataset

	ProPublica Full Data	ProPublica 2-Yr Data
0	6858	3743
1	3473	3471
Total	10331	7214

If we look at a Table with recidivism status by the pre versus post-April 1, 2014 COMPAS screen date indicator, we see that ProPublica incorrectly kept almost one thousand extra recidivists in the two-year general recidivism dataset.

Table 2: Any Recidivism by Pre-Post April 1 2014 COMPAS screen date - ProPublica Two-Year Dataset

is_recid	post_april_2014		
	0	1	Total
0	3743	0	3743
1	2473	998	3471
Total	6216	998	7214

The 998 recidivists who were incorrectly kept in the two-year (general recidivism) data represent a 28.8% share of the 3,471 recidivists in that dataset. Alternatively, we can say that ProPublica kept 998/2473 or 40.4% more recidivists than it should have. (And these shares are even higher for the slightly smaller subset of two-year recidivists; see Table 5 in the Recidivism Rates Section below)

I construct a **corrected** version of the two-year general recidivism dataset where I simply drop all people with a COMPAS screen date after April 1, 2014, including recidivists.[18] In this corrected dataset, I end up with the same number of non-recidivists as in the ProPublica two-year dataset, but I have 998 fewer recidivists. If we look at the COMPAS screening dates for this corrected dataset, we have the following:[19]

---

[17]The difference of 2 people is because these 2 people have a value of "N/A" for the COMPAS score category, and ProPublica drops these from the two-year recidivism datasets.

[18]To avoid right-censoring due to jail and/or prison time served for the original offense, one should perhaps implement an even earlier COMPAS screen date sample cutoff. Although ProPublica already dealt with this issue (for non-recidivists). I discuss the optimal cutoff further in the Appendix. In any case, the results presented here with the April 1, 2014 cutoff are robust to using an earlier (optimal) cutoff instead. Therefore, for simplicity in exposition and comparability to ProPublica's two-year dataset, I use the April 1, 2014 cutoff.

[19]Again, for comparison to the full data Figures displayed earlier, I do this for the overall or any recidivism variable (i.e. the "is_recid" variable in ProPublica's dataset; instead of the "two_year_recid" variable).
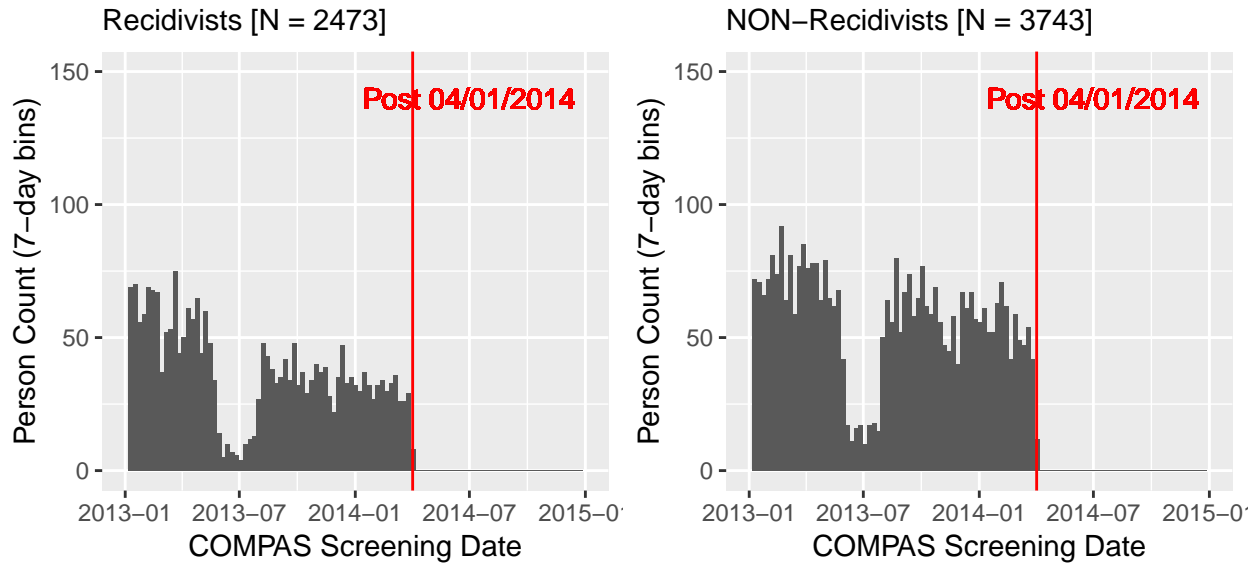
Figure 5: Persons by COMPAS Screen Date (7-day bins) by Recidivism Status - Corrected Two-Year Dataset

The graph on the right-side panel, for non-recidivists, is exactly the same as in the previous Figure. However, the graph on the left-side panel, for recidivists, now correctly drops everyone with a COMPAS screen date post-April 1, 2014.

4   Recidivism Rates

Here I will focus on the two-year recidivism variable. The COMPAS screen date Figures above, split by recidivism status, were done using the overall or any recidivism variable "is_recid", not the "two_year_recid" variable, for comparison purposes to the full data. As we see here, there are 220 more people with is_recid=1 than two_year_recid=1. These are people who recidivated, but did so after more than two years after the original COMPAS screen date (but before the end date of ProPublica criminal history data window at the end of March 2016).

Table 3: Any vs. Two-Year Recidivism - ProPublica Two-Year Dataset

is_recid	two_year_recid		
	0	1	Total
0	3743	0	3743
1	220	3251	3471

These 220 people represent only a 0.06 share of the 3,471 people who recidivate in total. (These 220 people, by definition, all have COMPAS dates before April 1, 2014) So the main findings of the data processing error are similar for the overall or any recidivism variable and the two-year recidivism variable. Either way, the bias in ProPublica's two-year recidivism datasets is clear: there is a disproportionate number of recidivists.

This fundamental problem in dataset construction affects some statistics more than others. It obviously

has a substantial impact on the total number of recidivists, and hence, also the share or rate of recidivism. In particular, it *artificially inflates* the recidivism rate. The artificially inflated two-year recidivism rate in ProPublica's dataset is the following:

Table 4: Two-Year Recidivism Rate - ProPublica Two-Year Dataset

two_year_recid	Two-Year Recidivism	
	N	Rate
0	3963	0.549
1	3251	0.451

If we repeat the earlier Table 2, which displays recidivism status by the post-April 1, 2014 COMPAS screen indicator, but now using the two-year recidivism indicator instead of the overall or any recidivism indicator (i.e. two_year_recid instead of is_recid), we have the following Table:

Table 5: Two-Year Recidivism vs. Pre-Post April 1 2014 COMPAS screen date - Two-Year Data

two_year_recid	post_april_2014		
	0	1	Total
0	3963	0	3963
1	2253	998	3251
Total	6216	998	7214

The 998 recidivists who were incorrectly kept in ProPublica's two-year data represent a 30.7% share of the 3,251 people who recidivated within two years in that dataset.[20] Alternatively, we can say that ProPublica kept 998/2253 or 44.3% more two-year recidivists than it should have.

Table 6: Two-Year Recidivism Rate vs. Pre-Post April 1 2014 COMPAS screen date - Two-Year Data

two_year_recid	post_april_2014	
	0	1
0	0.638	0.000
1	0.362	1.000

From these Tables, we see that the two-year recidivism rate is 45.1% in ProPublica's two-year data. We also see again that since ProPublica kept recidivists (but did not keep non-recidivists) with COMPAS screen dates post 4/1/14, all people with COMPAS screen dates post 4/1/2014 in the two-year recidivism dataset are recidivists. Therefore, as shown on the Table the correct two-year recidivism rate for the two-year data should be 36.2%. But due to the sizable group of post 4/1/2014 recidivists that ProPublica incorrectly kept, the two-year recidivism rate is artificially inflated to 45.1%. So there is a difference of 8.8 percentage points,[21] and thus, the two-year recidivism rate calculated by ProPublica is 24.3% higher than the true rate.

To test whether the difference in the recidivism rates is statistically significant, I compare the rate obtained with ProPublica's two-year dataset against the rate in the corrected two-year dataset, and vice-versa. I

---

[20]This 998 total is the same as before since by definition all the people with is_recid = 1 who have COMPAS screen dates after April 1, 2014, also have two_year_recid = 1 (since April 1, 2014, is less than two years before the end of the data window in late March 2016).

[21]While the difference of the rounded recidivism percent rates is 8.9 percentage points, with less rounding the respective rates are 45.07% and 36.25%, and hence, the actual difference is 8.82 percentage points, which is rounded to 8.8 percentage points.

use one-sample tests for this since these two datasets (and the statistics calculated from them) are not independent samples. I do two types of tests, a t-test, and a chi-squared test, which is more appropriate for comparing proportions or rates. (The null hypothesis, H_0, in each case, is the mean recidivism rate in the other dataset)[22]

Table 7: Statistical Significance Tests: Recidivism Rate - ProPublica vs. Corrected Two-Year Datasets

	N	Mean	SE	Low CI	Hi CI	Null Ho	t-stat	p-val.	chi-sq.	p-val.
ProPub_vs_Correct	7214	0.451	0.006	0.439	0.462	0.362	15.06	2e-50	242.9	9e-55
Correct_vs_ProPub	6216	0.362	0.006	0.35	0.374	0.451	-14.46	1e-46	195.3	2e-44

Given the small standard errors of 0.006, the difference in the recidivism rate (or the mean recidivism) between the two datasets, which is 8.8 percentage points, is highly statistically significant (p-values are very small, and hence, are displayed in scientific notation).

One can further examine the difference in the recidivism rates between the two datasets for people across the COMPAS score (decile) distribution. I do so in the next Figure.

Figure 6: Two-Year Recidivism Rate by COMPAS score decile
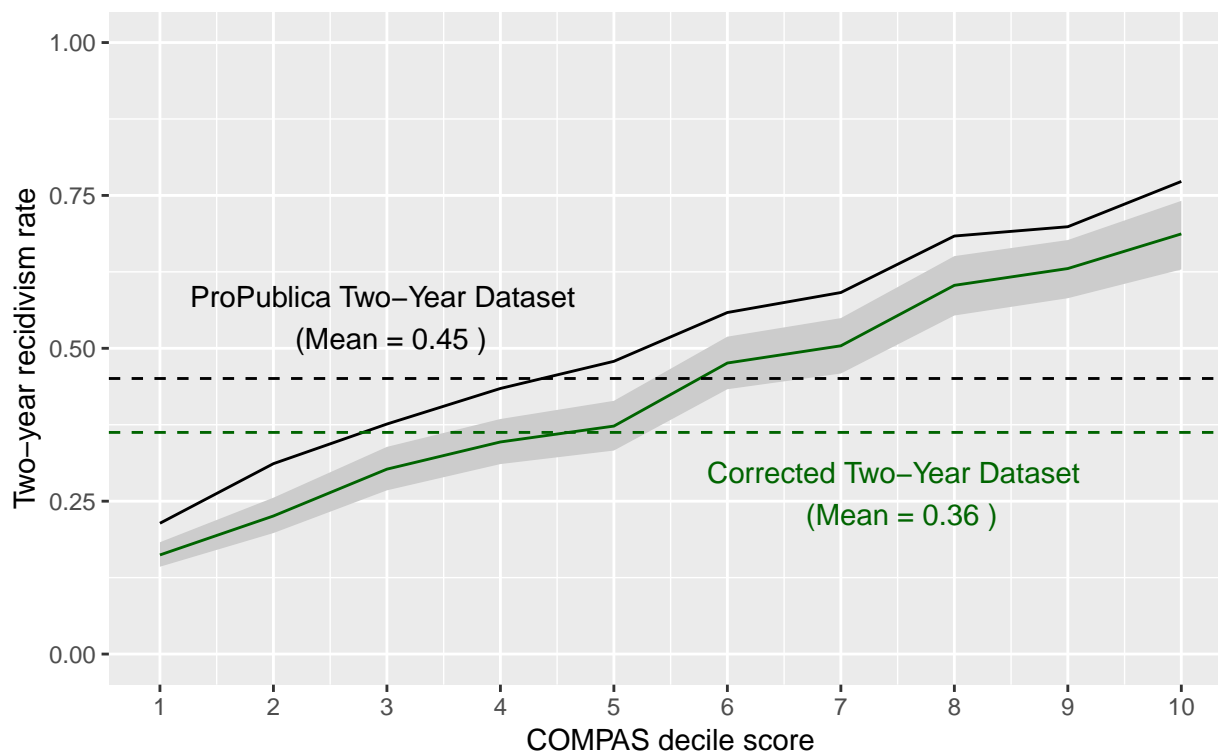
As we see in this Figure, the two-year recidivism rate is clearly higher for the ProPublica two-year dataset (black curve) than the corrected dataset (dark green curve) at every COMPAS score decile. Moreover, this

---

[22]I report results for two-sided tests, although one could in principle do one-sided tests here. Those would be even more statistically significant.

difference is statistically significant at every decile (I show the confidence interval for the corrected two-year data only since this dataset and the ProPublica two-year dataset are not independent datasets).

Another way of seeing ProPublica's data processing mistake when creating the two-year recidivism datasets is by doing a survival analysis. In the Appendix, I do such an analysis, and it confirms the results presented here.

5 Confusion Matrix / Truth Tables

Here I explore how the data processing mistake impacts other results. In particular, I look at the effect on the results from the contingency table analysis performed by ProPublica. For this analysis, ProPublica turned the COMPAS score categories of Low, Medium, and High, into a binary classifier, grouping Medium and High scores into an overall High score category.[23] I do the same here and report the results obtained by ProPublica with its two-year recidivism dataset, and the analogous set of results obtained using the corrected two-year recidivism dataset.

Table 8: ProPublica Two-Year Dataset: COMPAS Score Categories

factor_score_text	Freq
Low	3897
Medium	1914
High	1403
Total	7214

Table 9: ProPublica Two-Year Dataset: COMPAS Score Categories (Converted to Binary)

high_score	Freq
0	3897
1	3317
Total	7214

Table 10: ProPublica Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual	Predicted		
two_year_recid	COMPAS Score		
	Low	High	
0	2681	1282	0.55
1	1216	2035	0.45

---

[23]The initial mapping by Northpointe of COMPAS score *deciles* to Low, Medium, and High is score deciles 1-4, 5-7, and 8-10, respectively. Thus, for the further collapsed binary score where one groups Medium and High scores into an overall High score category, the decile mapping for Low vs. High scores is deciles 1-4 vs. 5-10.

Table 11: Corrected Two-Year Dataset: COMPAS Score Categories (Converted to Binary)

high_score	Freq
0	3522
1	2694
Total	6216

Table 12: Corrected Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual two_year_recid	Predicted COMPAS Score		
	Low	High	
0	2681	1282	0.64
1	841	1412	0.36

Table 13: Confusion Matrix Results that are similar between ProPublica vs. Corrected Two-Year Datasets

		Accuracy	FPR	FNR
ProPublica_results	7214	0.654	0.323	0.374
Corrected_results	6216	0.658	0.323	0.373

Table 14: Confusion Matrix Results that are different between ProPublica vs. Corrected Two-Year Datasets

		Prevalence	Pos Pred Value	Neg Pred Value	Detection Rate
ProPublica_results	7214	0.451	0.61	0.69	0.28
Corrected_results	6216	0.362	0.52	0.76	0.23

In the Tables above, I have replicated some of the results obtained by ProPublica.[24] I also report the analogous results using the corrected sample cutoff two-year recidivism dataset. In addition to the prevalence of recidivism (i.e. the recidivism rate), which we already discussed in the previous Recidivism Rates Section, we see that the biased two-year dataset used by ProPublica also affects the positive predictive value (PPV) (which is often referred to as "precision"), the negative predictive value (NPV), and the detection rate. In the corrected data, with the lower prevalence of recidivism, not surprisingly, we see in the Table above that PPV (and the detection rate) is lower and NPV is higher. Northpointe focuses on the complements to PPV and NPV; i.e. 1 minus these (see Dieterich et al. 2016). If we focus on these complements instead, we see that with the biased ProPublica two-year dataset a 0.39 (i.e. 1 - 0.61) share of people was labeled high risk but did not re-offend, whereas with the corrected data a 0.48 share of people is labeled high risk but does not re-offend. And before a 0.31 share was labeled low risk but did re-offend, whereas now a 0.24 share does so. Again, given the lower prevalence of recidivism in the corrected data, it is not surprising that one type of error goes up and the other goes down.

On the other hand, as we also see in the Tables above, the biased dataset has relatively little impact on several other key statistics, such as accuracy, the false positive rate (FPR), and the false negative rate (FNR).[25] This

---

[24]In particular, those reported in item (51) in their GitHub Jupyter notebook (Larson et al. 2017). Although the accuracy and detection rates are not reported by ProPublica.

[25]Or one minus these rates, i.e. specificity and sensitivity.

is perhaps not that surprising. The FPR is by definition independent of the actual number of positives (or recidivists) in the data since it is the ratio of the number of cases *predicted* to be positive (or to recidivate) but that are *not* actually positive, over all the cases that are *not* positive. So the FPR is calculated based only on actual negatives (i.e. only people who do not recidivate). As a result, the FPR in the corrected two-year data is *exactly* the same as the FPR in ProPublica's two-year data.

At the same time, the FNR is only based on actual positives or people who do recidivate. The FNR is the ratio of people who are *predicted* not to recidivate but who actually recidivate, over all people who recidivate. The total count of people who recidivate is clearly quite different between the two datasets. However, as long as the COMPAS scores of the additional recidivists who ProPublica incorrectly kept in the two-year data is similar to the COMPAS scores of the recidivists who are correctly kept in the two-year data, then this will have little effect on the FNR. We see that this is indeed the case in the following Table:

Table 15: COMPAS Low/High Score vs. Pre-Post April 1 2014 COMPAS screen date - Two-Year Data

	Recidivists Only	
High COMPAS score	post_april_2014	
	0	1
0	841	375
1	1412	623
Total	2253	998

Table 16: COMPAS Low/High Score vs. Pre-Post April 1 2014 COMPAS screen date - Two-Year Data

	Recidivists Only	
High COMPAS score	post_april_2014	
	0	1
0	0.37	0.38
1	0.63	0.62

In this Table we see that the Low-High COMPAS score distribution is almost identical for recidivists with COMPAS screen dates prior to the two-year cutoff date of April 1, 2014, as it is for recidivists with COMPAS screen dates after that date; the share of recidivists with a high COMPAS score is 63% for the former, and 62% for the latter. Therefore, it is not surprising that the FNR for the corrected two-year dataset is almost identical to the FNR for ProPublica's original two-year dataset.[26]

Next, following ProPublica's analysis, I repeat the confusion matrix analysis separately for African-Americans and Caucasians (whom I label blacks and whites, respectively, in the Tables below). This is the key analysis that garnered the most attention when ProPublica's article was published, with a higher false positive rate (FPV) and a lower false negative rate (FNR) for blacks than whites. (I just show the results here; the actual confusion matrix tables themselves are in the Appendix)

---

[26]This similarity in the FNRs, combined with identical FPRs as we saw previously, means that there is also very little impact on the receiver-operating characteristic (ROC) curve and the area under that curve; see Appendix.

Table 17: Blacks: Confusion Matrix Results that are similar between ProPublica vs Corrected Two-Year Data

		Accuracy	FPR	FNR
ProPublica_results_blacks	3696	0.638	0.448	0.280
Corrected_results_blacks	3139	0.624	0.448	0.279

Table 18: Whites: Confusion Matrix Results that are similar between ProPublica v Corrected Two-Year Data

		Accuracy	FPR	FNR
ProPublica_results_whites	2454	0.670	0.235	0.477
Corrected_results_whites	2132	0.689	0.235	0.488

Table 19: Blacks: Confusion Matrix Results that do change between ProPublica vs Corrected Two-Year Data

		Prevalence	Pos Pred Value	Neg Pred Value	Detection Rate
ProPublica_results_blacks	3696	0.51	0.63	0.65	0.37
Corrected_results_blacks	3139	0.43	0.55	0.73	0.31

Table 20: Whites: Confusion Matrix Results that do change between ProPublica vs Corrected Two-Year Data

		Prevalence	Pos Pred Value	Neg Pred Value	Detection Rate
ProPublica_results_whites	2454	0.39	0.59	0.71	0.21
Corrected_results_whites	2132	0.30	0.49	0.78	0.15

As expected from the discussion and combined race sample results earlier, the FPR is identical, and the FNR is very similar, with the corrected data, so blacks have a substantially higher FPR and lower FNR than whites in the corrected data too. This key finding by ProPublica, therefore, does not change with the corrected data.[27]

The utility of focusing on the differences in the FPR and FNR across race groups, however, has been called into question by other researchers (see for example Corbett-Davies and Goel 2018). Moreover, just like we saw with the combined race sample, we see substantial differences in other statistics. In particular, regarding recidivism prevalence, PPV, NPV, and the detection rate.

6   Conclusion

While ProPublica's COMPAS score and recidivism data are used in an increasing number of studies to test various definitions and methodologies of algorithmic fairness, researchers have taken the data 'as is' to test their methodologies, but have not examined closely the data itself for data processing issues. This paper, instead of testing a novel fairness definition or procedure, takes a closer look at the actual datasets put

---

[27]Here again, note that accuracy (and the detection rate) was not reported by ProPublica. The lack of reporting for accuracy, especially in these by-race results, is one of Northpointe's main critiques of ProPublica's analysis since the accuracy is similar for blacks and whites (Dieterich et al. 2016).

together by ProPublica. In particular, the sub-datasets built to study the likelihood of recidivism within two years of the original offense arrest and COMPAS screening date.

I take a new yet simple approach to visualize these data, by analyzing the distribution of defendants across COMPAS screening dates. Doing so, I find that ProPublica made an important data processing mistake creating these key datasets often used by other researchers. As I show in this paper, ProPublica failed to implement a two-year sample cutoff rule for recidivists (whereas it implemented an appropriate two-year sample cutoff rule for non-recidivists). As a result, the bias in the two-year dataset is clear, there are a disproportionate number of recidivists. To my knowledge, this is the first paper to highlight this key data processing mistake.

When I implement a simple two-year COMPAS screen date cutoff rule for all people, including recidivists, I estimate that in the two-year general recidivism dataset ProPublica kept 44.3% more two-year recidivists than it should have. This fundamental problem in dataset construction affects some statistics more than others. It obviously has a substantial impact on the recidivism rate. In particular, it artificially inflates the two-year general recidivism rate by 8.8 percentage points, from 36.2% to 45.1%, which represents a 24.3% increase in the two-year recidivism rate.

ProPublica's data processing mistake also affects the positive predictive value (PPV) or precision, and the negative predictive value (NPV). On the other hand, it has relatively little impact on several other key statistics that are less susceptible to changes in the relative share of recidivists versus non-recidivists, such as accuracy, the false positive rate (FPR), and the false negative rate (FNR). While the latter statistics, especially the differentials in the FPR and the FNR by race, have garnered the most attention in the academic research and public debate, the utility of focusing on those particular metrics has been called into question by other researchers (see Corbett-Davies and Goel 2018).

Ultimately, the practical importance of this data processing mistake may be limited. I am not suggesting that Northpointe itself made a mistake in actually developing the COMPAS recidivism risk score. While the data used for that, and the actual model, are proprietary and not publicly available; it is unlikely that a similar mistake was made when developing such scores, or other recidivism risk scores by other companies.[28] Although domain expertise does not always translate into correctly processed data. For example, Northpointe's critique of ProPublica's analysis, using ProPublica's COMPAS datasets (Dieterich et al. 2016), fails to identify ProPublica's data processing mistake, and thus, produces some biased results. Or the rejoinder to ProPublica by two criminal justice academics and a judicial system administrative officer, who also do not identify the mistake in ProPublica's data and hence produce Figures where the two-year recidivism rate is biased upward (Flores et al. 2016). In any event, this paper puts the focus on, and highlights the potential pitfalls in, the data processing stage. I am currently working on a GitHub repository to make public the corrected data, although the data correction is straightforward and can be implemented by others independently.[29]

References

Angwin, J., Larson, J., Mattu, S., and Kirchner, L. (2016), "Machine Bias. There's software used across the country to predict future criminals. And it's biased against blacks." *ProPublica.*

Barocas, S., and Selbst, A. D. (2016), "Big data's disparate impact," *California Law Review*, HeinOnline, 104, 671.

Bilal Zafar, M., Valera, I., Gomez Rodriguez, M., and Gummadi, K. P. (2016), "Fairness Beyond Disparate

---

[28]Moreover, it is not clear to what extent a recidivism risk score is used by judges at the *pretrial* stage to set bail. Although Cowgill using the ProPublica COMPAS data finds a non-trivial effect at score class breakpoints (Cowgill 2018). (I am not sure if Cowgill corrected the data processing issue highlighted in this paper when doing his analysis, and if doing so would have any impact on his results)

[29]Additionally, Rudin et al. (2018) have also reconstructed the ProPublica COMPAS datasets from the original ProPublica Python database and made them available on GitHub. In so doing they may have avoided making the same data processing mistake as ProPublica. (Although they do not generally highlight the dataset differences between their dataset and ProPublica's, and do not identify ProPublica's data processing mistake. As mentioned earlier, their focus is altogether different)

Treatment & Disparate Impact: Learning Classification without Disparate Mistreatment," *arXiv e-prints*, arXiv:1610.08452.

Bilal Zafar, M., Valera, I., Gomez Rodriguez, M., Gummadi, K. P., and Weller, A. (2017), "From Parity to Preference-based Notions of Fairness in Classification," *arXiv e-prints*, arXiv:1707.00010.

Chouldechova, A. (2016), "Fair prediction with disparate impact: A study of bias in recidivism prediction instruments," *arXiv e-prints*, arXiv:1610.07524.

Corbett-Davies, S., and Goel, S. (2018), "The Measure and Mismeasure of Fairness: A Critical Review of Fair Machine Learning," *arXiv e-prints*, arXiv:1808.00023.

Corbett-Davies, S., Pierson, E., Feller, A., Goel, S., and Huq, A. (2017), "Algorithmic decision making and the cost of fairness," *arXiv e-prints*, arXiv:1701.08230.

Cowgill, B. (2018), "The Impact of Algorithms on Judicial Discretion: Evidence from Regression Discontinuities," *Working Paper*.

Cowgill, B., and Tucker, C. E. (2019), "Economics, fairness and algorithmic bias," *SSRN Electronic Journal*.

Dieterich, W., Mendoza, C., and Brennan, T. (2016), "COMPAS Risk Scales: Demonstrating Accuracy Equity and Predictive Parity," *Northpointe Inc.*

Flores, A. W., Bechtel, K., and Lowenkamp, C. T. (2016), "False Positives, False Negatives, and False Analyses: A Rejoinder to "Machine Bias: There's Software Used Across the Country to Predict Future Criminals. And It's Biased Against Blacks.","" *Federal Probation Journal*, 80 Number 2.

Kleinberg, J., Ludwig, J., Mullainathan, S., and Rambachan, A. (2018), "Algorithmic fairness," *AEA Papers and Proceedings*, 108, 22–27.

Larson, J., Mattu, S., Kirchner, L., and Angwin, J. (2016), "How We Analyzed the COMPAS Recidivism Algorithm," *ProPublica*.

Larson, J., Mattu, S., Kirchner, L., and Angwin, J. (2017), "COMPAS Analysis.jpynb," *ProPublica Jupyter Notebook on GitHub*.

Rudin, C., Wang, C., and Coker, B. (2018), "The age of secrecy and unfairness in recidivism prediction," *arXiv e-prints*, arXiv:1811.00731.

7 Appendix

7.1 Full Dataset Drops

When creating the two-year recidivism datasets, ProPublica reduced the full data sample for two reasons even before implementing its (faulty) two-year cutoff. If we take the full dataset as a starting point, with 11,757 people, ProPublica for some reason dropped the last 756 person IDs when constructing the two-year datasets. Starting from person ID 11002 to person ID 11757 in the full dataset. It is not clear why these people were dropped. Many have COMPAS screen dates prior to 4/1/2014 since person IDs are not chronologically ordered. And thus, many of these people are observed for two years or recidivate within two years. In any case, I also dropped these 756 people in the construction of the corrected two-year recidivism dataset, to make it as comparable as possible to ProPublica's 7,214 person two-year general recidivism dataset (but for the explicit two-year COMPAS screen date cutoff correction I implement).

Additionally, ProPublica also dropped 719 people who did not appear to have good data. ProPublica could not find case/arrest information on these people. ProPublica tagged these as "is_recid = -1" in the full dataset.[30] There is some overlap between the 756 people mentioned in the previous paragraph and these

---

[30]Interestingly, ProPublica dropped these people from the two-year general recidivism dataset. But it did not drop them from the two-year *violent* recidivism dataset. (While it did drop them from the more reduced 4,743 two-year violent csv file, it did not drop them in the 6,454 two-year violent data it used for the violent recidivism truth tables)

719 people. So the net additional drop in this step is actually 670 people. Thus, one ends up with 10,331 people total in the 'full' dataset.[31] I also drop these 719 (or 670 *additional*) people in the construction of the corrected two-year general recidivism dataset, so as to make it more comparable to ProPublica's two-year general recidivism dataset (again, but for the explicit two-year COMPAS screen date cutoff correction).

7.2 Assumptions Regarding Data and Analysis

This paper's key objective has been to point out the fundamental data processing error made by ProPublica in the construction of its two-year recidivism datasets. As such, I do not engage in a wholesale revision of the ProPublica data and analysis. Therefore, I mostly take as given many aspects of the data and analysis, and make many of the same assumptions made by ProPublica and other researchers. (While I may revisit some of these assumptions in future work, that is not the purpose of the current paper) Therefore, I am otherwise assuming the data is generally in good shape, and that the analytic approach is valid. However, here I list some exceptions to this assumption regarding the quality of the data. I also list the key assumptions made in the analysis.

Data Assumptions

- As with many data collection efforts that must obtain different features on a given sample from different data sources, and then match these, the matching is not perfect, and ProPublica acknowledges this:

  "We found that sometimes people's names or dates of birth were incorrectly entered in some records – which led to incorrect matches between an individual's COMPAS score and his or her criminal records. We attempted to determine how many records were affected. In a random sample of 400 cases, we found an error rate of 3.75 percent (CI: +/- 1.8 percent)." (Larson et al. 2016)

  I have not explored this data matching issue in my analysis.

- Related to this matching issue, there are some people in its data who have multiple COMPAS screen dates. In calculations not shown, I find there appear to be 688 people in the 11,757 person dataset with multiple COMPAS screen dates. ProPublica seems to have selected a single COMPAS screen date for such people when it builds the two-year dataset(s) (as well as its Cox datasets). I have not explored how it selected a single date. But since it is a relatively small share of people who have multiple COMPAS screen dates to begin with, this should not affect the main findings in my paper.

- Also related to this, there are some people who ProPublica finds do not have good data. In particular, ProPublica says it could not find some key case and/or arrest information for these people. They total 719 out of the 11,757 people in the full dataset or 6.1%. I also drop these people.[32]

- A very small number of people appear to have implausible *negative* time spells outside prison. In calculations not shown here, I find that in the 11,757 person full dataset, only 63 people have such negative time spells. ProPublica adds these negative amounts (as negative) when calculating the total time outside of prison for such people. I do the same.

- Some people have a "current" offense date that occurs a long time prior to the COMPAS screen date. However, the 'jail_in' date for these people is close to the COMPAS screen date, so such people could plausibly have committed the offense a long time ago and only been caught/charged recently. So they do not necessarily represent a data problem.

- ProPublica obtained criminal history information from the Broward County Clerk's Office website, and jail records from the Broward County Sheriff's Office, as well as public incarceration records from the Florida Department of Corrections website. I am not sure what happens if someone in ProPublica's data sample moves away from Florida after the COMPAS screen date. In particular, it is not clear

---

[31]This is also the same number of people as in ProPublica's Cox general recidivism dataset "cox-parsed.csv".

[32]As discussed in the first section of this Appendix, ProPublica tagged these as "is_recid = -1" in the full dataset.

whether that person would show up in the data again if he or she commits a crime in a different state. I also do not know what happens if any of the people in the sample become deceased. There could be some sample attrition.

- There are two months with very few people with COMPAS screen dates (June and July 2013).[33] It is not clear why there is such a drop in COMPAS cases during these two months. To the extent this is a problem, it appears to be a problem with the original dataset that ProPublica received from Broward County since it is also evident in the "compas-scores-raw.csv" dataset. So it does not appear to be a data processing issue by ProPublica. Thus, it is not clear what can be done about this (unless perhaps if one goes back to the original source in Broward County to try to collect the data again). In any event, I checked whether the relatively few people with COMPAS screen dates during these two months looked different in various dimensions, but they did not. (Except they did have a slightly longer period of time between the 'jail_in' date and COMPAS screen date, with a mean of 3.3 days, compared to 0.6 days for the rest of the data)

- As I discuss in the first section of this Appendix, for some reason ProPublica dropped the people with the last 756 person IDs in its original pretrial defendants sample. It is not clear why it dropped these people. However, I also drop them for comparability to ProPublica's analysis.

- As other researchers note, some people in this dataset have low COMPAS scores and yet, surprisingly, have many prior offenses (Rudin et al. 2018). These researchers also note that on the flip-side, some people have high COMPAS scores, but no priors, and their current offense is non-violent (for this group, the researchers hypothesize that maybe ProPublica's data is missing some criminal history information).

- Finally, the age variable that ProPublica constructed is not quite accurate. ProPublica calculated age as the difference in years between the point in time when it collected the data, in early April 2016, and the person's date of birth. However, when studying recidivism, one should really use the age of the person at the time of the COMPAS screen date that starts the two-year time window. So some people actually may be up to two years younger than the age variable that ProPublica created. Since I do not really use age in any of my analyses, I do not take the trouble of correcting this variable.

Analysis Assumptions

- Since this analysis is for people in Broward County and for a particular point in time, it may not generalize to other jurisdictions and time windows.

- In this paper, I have focused only on the two-year general recidivism dataset, and not the smaller dataset(s) that ProPublica created for the sub-category of two-year violent recidivism, although the latter suffers from the same data processing mistake as the two-year general recidivism dataset.[34]

- The observed recidivism rate is really a re-arrest rate. It may not reflect the true recidivism rate in the sense that some people may commit new offenses but not get caught. (Clearly, therefore, the amount and aggressiveness of policing may affect the observed recidivism rate)

- I focus on the study of the fixed time-period two-year recidivism outcome. With survival data, however, it is often preferable to apply survival models. Although, in the survival analysis section later in this Appendix, I show that at the two-year mark, the two approaches are almost identical (at least without controls). A survival analysis, nonetheless, gives a fuller picture of recidivism, since it is not constrained to a single point in time.

- Foregoing the fuller picture provided by a survival analysis approach, for the most part, and doing an analysis of recidivism at a particular point in time instead, I am assuming that the two-year recidivism

---

[33]Also noticeable is the higher number of COMPAS screen dates in the first half of 2013.

[34]ProPublica did not actually make readily available the two-year violent recidivism csv file it uses for its key violent recidivism analysis. But it can be reconstructed. The even more reduced two-year violent recidivism csv file ProPublica did make available, which it uses only in certain parts of its analyses, has the further problem that it drops people who are non-violent recidivists entirely (instead of tagging them as non-recidivists for violent offenses).

metric is the appropriate recidivism metric for this approach. (As opposed to, say, one-year recidivism, or three-year recidivism, etc.) ProPublica explains why it chose this particular time-frame, saying it:

> "based this decision on Northpointe's practitioners guide, which says that its recidivism score is meant to predict 'a new misdemeanor or felony offense within two years of the COMPAS administration date.' "[35]

- I am assuming that netting out prison (and jail) time served for the original offense, and only keeping non-recidivists who are observed for more than two-years outside of jail/prison is appropriate. This sub-setting seems reasonable, since clearly the recidivism rate may be quite different in prison than outside of prison.

  For consistency, one should also drop recidivists who only recidivate while in prison for their original offense. However, it does not appear that one can identify such recidivists readily from a feature in the data. Nevertheless, one could in principle estimate whether the recidivism offense date occurs during a time window when the person is in custody. Exploring this (in calculations not shown), I find that less than 2% of recidivists seem to commit a recidivism offense while in custody for their original offense. If one were to drop these recidivists, the two-year recidivism rate would drop further, but by less than 1 percentage point.

  Alternatively, to avoid most of the right-censoring due to prison time for the original offense, starting from the full dataset, one could potentially just keep all non-recidivists and recidivists whose COMPAS screen dates are prior to an "optimal" COMPAS screen cutoff date that is a non-trivial amount of time prior to the April 1, 2014 cutoff, regardless of how much time they subsequently spend in jail or prison. (But still net out prison/jail time to construct the two-year recidivism indicator flag) This is potentially useful in order to have both recidivists and non-recidivists on a more equal footing. I explore this issue further below.

- In the analyses that utilize the COMPAS score, I am assuming that it is valid to study the COMPAS recidivism risk score for *pretrial* defendants. As Flores et al. (2016) point out, the recidivism risk score may actually be intended to be applied more to current prison inmates for probation decisions. (Indeed, the ProPublica data has a different set of COMPAS scores, regarding the risk of failure to appear in court, which may be intended for pretrial decisions instead. I have not explored this alternative COMPAS score)

- I do not explore any feedback loop effects. As Cowgill points out and examines, judges sometimes use COMPAS scores to guide their *bail* decisions "and longer bailtime exerts a causal influence on defendants' outcomes, including recidivism." (Cowgill 2018)

- The contingency table analyses assume that using a binary score category (Low vs. High) for the predictor variable is adequate. As opposed to a more detailed score breakdown, such as Low, Medium, High, or score deciles, or the continuous raw score. And that the breakpoint used, which groups deciles 1-4 and 5-10 into the two categories is appropriate (although I briefly examine different thresholds for the binary split of score deciles; see the ROC Curves Section below).

7.3 Survival Analysis

Another way of seeing ProPublica's data processing mistake when making the two-year recidivism datasets is by doing a survival analysis. In the Figure below I graph the Kaplan-Meier survival curves for the full data, the ProPublica two-year general recidivism data, and the corrected two-year data. I use the overall or any recidivism variable "is_recid", not the "two_year_recid" variable, here, so we can see the full curve, even past two years.[36]

---

[35]ProPublica also points to "a (recent) study of 25,000 federal prisoners' recidivism rates by the U.S. Sentencing Commission, which shows that most recidivists commit a new crime within the first two years after release (if they are going to commit a crime at all)."

[36]As we saw in the Recidivism Rates Section above, in the two-year datasets there are 220 more people with is_recid=1 than two_year_recid=1. These are people who recidivated, but did so more than two years after the original COMPAS date (but
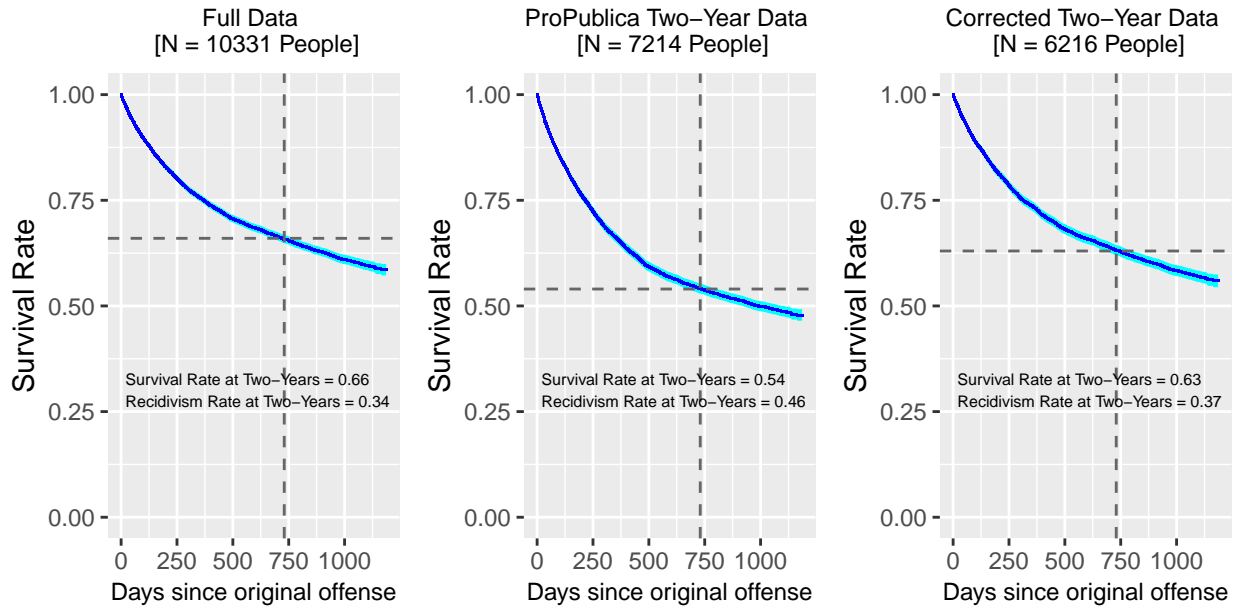
Non–Recidivism Survival Curve

Figure 7: Non-Recidivism Survival Curves - Three Datasets

As we see from these graphs, at the two-year mark a 0.34 share of people has recidivated in the full data. However, in ProPublica's two-year general recidivism data, at the two-year mark, a much higher fraction of people recidivates after two years, 0.46. (This rate is almost identical to the rate estimated in the Recidivism Rates Section above, 0.45). In the corrected two-year data, at the two-year mark, a 0.37 share of the people has recidivated, which is much closer to the full data estimate of 0.34.[37] (And is also almost identical to the rate estimated for the corrected two-year data in the Recidivism Rates Section above, 0.36).[38]

7.4 Optimal COMPAS Screening Date Cutoff

Turning back now to the fixed time window two-year analysis, to avoid right-censoring for non-recidivists due to prison time served for the offense they committed just prior to their COMPAS screening (i.e. the offense that led them to be jailed and screened, and perhaps later convicted and imprisoned), one should really implement a COMPAS screen sample cutoff date *earlier* than April 1, 2014. Indeed, ProPublica already deals with this for non-recidivists in the two-year data, since it nets out any jail/prison time they served for the original offense, and only keeps non-recidivists observed for two years outside of jail/prison. This is somewhat analogous to (but not exactly the same as) implementing an earlier COMPAS screen-date cutoff (for non-recidivists). (Of course, ProPublica implements no cutoff at all for recidivists) Since ProPublica, therefore, drops non-recidivists with considerable jail and/or prison time for their original offense, one should really implement an overall COMPAS screen date cutoff (i.e. for recidivists too) that is actually some non-trivial amount of time prior to April 1, 2014, to put recidivists and non-recidivists on a more equal footing.[39]

---

before the end date of ProPublica's criminal history data window at the end of March 2016).

[37]The slight difference between the corrected two-year data and the full data is due to the sample composition difference. The corrected two-year data does not contain any people with COMPAS dates post-April 1, 2014, so we shouldn't expect the rates to be exactly the same.

[38]If we use three digits after the decimal point, the recidivism rates estimated in the survival analysis are as follows, 0.342, 0.459, and 0.368, for the full dataset, the ProPublica two-year dataset, and the corrected two-year dataset, respectively. For comparison, in the Recidivism Rates Section earlier, the last two rates were 0.451 and 0.362, respectively. So the survival analysis results are almost identical to the results in that earlier section.

[39]As noted in the assumptions section earlier in this Appendix, to really treat them equally, one should further drop any recidivists who recidivate while serving jail/prison time for their original offense; although this appears to be a small share of

This earlier COMPAS screen date cutoff is what I term the "optimal" cutoff. Using the full dataset of 10,331 defendants before ProPublica's two-year drops, here I explore what the optimal cutoff date might be for non-recidivists (but which would be applied across the board) that still preserves the most data. In the next Figure, focusing only on non-recidivists, I plot the fraction of people observed for two or more years outside jail or prison by COMPAS screen date.

Fraction of People by COMPAS screen date observed for 2+ years (outside jail or prison)
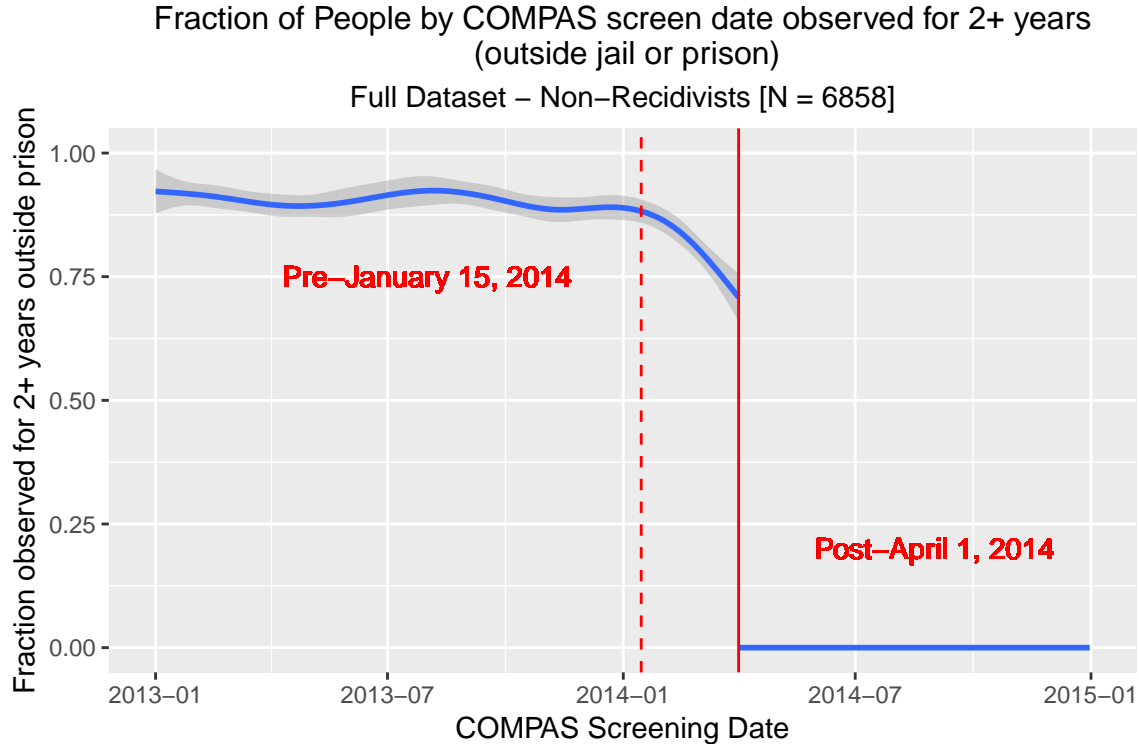
Full Dataset – Non–Recidivists [N = 6858]

Figure 8: Exploring Optimal COMPAS Screening Date Cutoff for Two-Year Recidivism Analysis

The solid red vertical line is, as always, on April 1, 2014. Clearly, as expected, after April 1, 2014, the fraction of people by COMPAS screen date observed for two or more years is zero. On the flip side, we see that prior to sometime in January 2014, the vast majority of non-recidivists are observed for two or more years outside jail or prison. I plotted a dashed red line on January 15, 2014, as reference. Between mid-January, 2014, and April 1, 2014, the fraction of non-recidivists observed for two or more years declines substantially (due to prison time served for the original offense) from approximately 88% to 71% (or 17 percentage points). However, prior to mid-January, this fraction is flatter and only rises from 88% to 92% (or only 4 percentage points). So the most appropriate sample cutoff that still preserves the most data appears to be sometime in early January, 2014.[40]

For simplicity in exposition (and comparability to ProPublica's two-year dataset), I have used the April 1, 2014 cutoff in most of this paper; which already pinpoints all the issues with ProPublica's two-year dataset. Moreover, the corrected recidivism rate is almost identical if one uses the April 1, 2014 cutoff, or an earlier January 2014 cutoff. If I use a January 15, 2014 COMPAS screen cutoff date, the two-year general recidivism rate is 36.3% instead of 36.2%. This is partly because any additional time that is cut off from ProPublica's two-year recidivism dataset relative to April 1, 2014, will drop both recidivists and non-recidivists (since ProPublica's data processing mistake only pertains to recidivists kept beyond April 1, 2014). To see this

---

recidivists.

[40]The smoothing method (function) used to create the non-parametric curve in this Figure is a generalized additive model (or 'gam'). The shape of the curve depicted here will vary somewhat depending on the smoothing method. So it is not straightforward to pinpoint an exact optimal COMPAS screen date cutoff point.

more clearly, in the next Table I display the two-year recidivism rates in the ProPublica two-year general recidivism dataset by COMPAS screen date grouped into year-quarter combinations.

Table 21: Two-Year Recidivism Rate vs. COMPAS screen date year-quarter - ProPublica Two-Year Data

two__year__recid	Year-Quarter							
	2013.1	2013.2	2013.3	2013.4	2014.1	2014.2	2014.3	2014.4
0	0.616	0.646	0.645	0.647	0.648	0.009	0.000	0.000
1	0.384	0.354	0.355	0.353	0.352	0.991	1.000	1.000

As we see from the Table above, an earlier cutoff will not affect in any meaningful way the results presented in the main body of this paper, since the two-year recidivism rate fluctuates in a relatively narrow band between 0.352 and 0.384 for people with COMPAS screen dates prior to April 1, 2014.[41] For most COMPAS screen date quarters the two-year recidivism rate actually fluctuates in a very narrow band between 0.352 and 0.355. In the first quarter of 2013, however, it is 0.384. It is not clear why the recidivism rate in this quarter is non-trivially higher, but the difference is not very large.[42] Moreover, given some likely underlying exogenous variation or random noise in the data, we would not necessarily have expected the recidivism rate to fluctuate in such an otherwise narrow range. Finally, we might have actually expected the recidivism rate to be slightly higher in the first quarter of *2014* since as just mentioned ProPublica already dropped from this two-year dataset non-recidivists with COMPAS screen dates prior to April 1, 2014, who have non-trivial jail and/or prison time, and hence, are not observed for a full two-years out of jail/prison. And this mostly affects non-recidivists in the first quarter of 2014 per the previous Figure. But in the first quarter of 2014, we observe a recidivism rate that is almost identical to previous quarters (and is actually *lower* than the first quarter of 2013). The fact that the recidivism rate for the first quarter in 2014 is not slightly higher may reflect some natural variability or noise in the data that is otherwise depressing and cancelling this anticipated effect.[43]

7.5   ROC Curves

With a classification variable that has several possible threshold values that one could use to turn it into a binary classifier, it is common practice to plot the receiver-operating characteristic (ROC) curve. The ROC plots the *sensitivity* (or 1 - FNR) against the *specificity* (or 1 - FPR) for a variable turned into a binary classifier, at various possible thresholds for the binary split. So far in this paper, I have focused on the binary score split that ProPublica used. ProPublica turned the COMPAS score categories of Low, Medium, and High, into a binary classifier, grouping Medium and High scores into an overall High score category. The initial mapping by Northpointe of the underlying COMPAS score *deciles* to the Low, Medium, and High score categories to begin with, is score deciles 1-4, 5-7, and 8-10, respectively. Thus, for the further collapsed binary score where one groups Medium and High scores into an overall High score category, the decile mapping for Low vs. High scores is deciles 1-4 vs. 5-10. However, one can explore all possible binary score breakpoints; e.g. score deciles 1 vs. 2-10, 1-2 vs. 3-10, etc.[44]

The way to explore this is by plotting the ROC curve, as I do here.[45] The area under the ROC curve (or the

---

[41] The recidivism rate in the second quarter of 2014, which starts on April 1, 2014, is just shy of 100% because ProPublica actually kept non-recidivists through the first day in that quarter, or through April 1, 2014. So one day in that quarter does have non-recidivists, which makes the recidivism rate in that second quarter of 2014 fall just short of 100% by approximately 1/120 (or 0.008) days.

[42] Recall also, as shown in Figures 1 and 2, that the COMPAS screen sample count is higher for the first quarter (and part of the second quarter) in 2013. So there may be some underlying differences in the sample across quarters.

[43] Indeed, in calculations not shown, I find that in the full dataset of 10,331 defendants *before* dropping non-recidivists with less than two years outside jail/prison, the recidivism rate in this quarter is about 2 percentage points lower than in each of the previous three quarters.

[44] Indeed, one could even use the raw score that is mapped into the score deciles, to begin with.

[45] ProPublica actually did this as well, separately for African-Americans and Caucasians; see https://github.com/propublica/compas-analysis/blob/master/Cox%20with%20interaction%20term%20and%20independent%20variables.ipynb.

AUC) gives a measure of how well the predictor variable performs in the aggregate across all possible binary thresholds that one can use to split it. The larger the area, the better the predictor (e.g. when comparing two different predictor variables, or alternatively, the same predictor on two different datasets). In the Figure below, I plot the ROC curves for the ProPublica versus corrected two-year datasets.
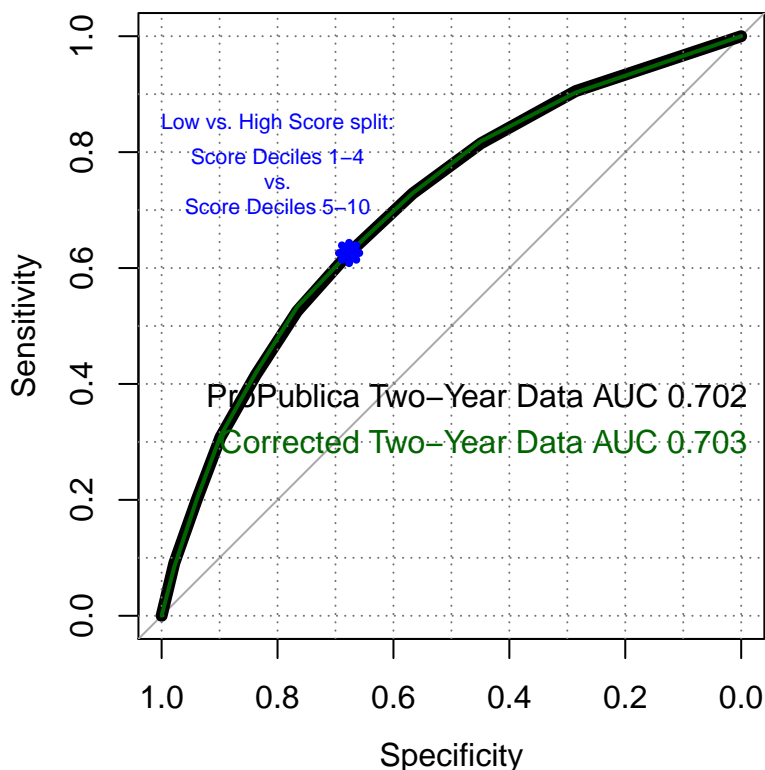

Figure 9: ROC curves for Two-Year Recidivism using COMPAS score deciles

As we can see from this Figure, the ROC curves for the two datasets are essentially identical, lying on top of one another (the ROC curve for ProPublica's two-year data is in black, and the ROC curve for the corrected data is in dark green). This is not that surprising given our earlier results in the Confusion Matrix Section, where we saw that the false positive rate is identical, and the false negative rate is very similar, across the two datasets (for the particular score decile split - deciles 1-4 vs. 5-10 - used there; which is indicated on this Figure by the blue dot mark). As discussed in that earlier section, the FPR is by definition independent of the actual number of positives (or recidivists) in the data; so it is identical in both datasets. At the same time, the FNR is only based on actual positives or people who do recidivate. While the total count of people who recidivate is clearly quite different between the two datasets, since the COMPAS scores of the *additional* recidivists who ProPublica incorrectly kept in the two-year data are similar to the COMPAS scores of the recidivists who are correctly kept in both datasets, then the FNR is very similar.

7.6 Confusion Matrix - Additional Results by Race

Focusing again on the score decile split analyzed in the main body of the paper, i.e. score deciles 1-4 vs. 5-10, at the end of the Confusion Matrix Section, we showed the confusion matrix analysis results separately by race. However, for ease of exposition, we did not show the actual confusion matrix contingency tables by race there. We do so here.

Table 22: Blacks ProPublica Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual two_year_recid	Predicted COMPAS Score		
	Low	High	
0	990	805	0.49
1	532	1369	0.51

Table 23: Whites ProPublica Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual two_year_recid	Predicted COMPAS Score		
	Low	High	
0	1139	349	0.61
1	461	505	0.39

Table 24: Blacks Corrected Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual two_year_recid	Predicted COMPAS Score		
	Low	High	
0	990	805	0.57
1	375	969	0.43

Table 25: Whites Corrected Two-Year Dataset Confusion Matrix: Recidivism vs. Low/High COMPAS Score

Actual two_year_recid	Predicted COMPAS Score		
	Low	High	
0	1139	349	0.7
1	314	330	0.3

7.7 Correcting Recidivism Figures in Other Papers

Finally, here I replicate some Figures in prior papers by other researchers who have used ProPublica's COMPAS two-year dataset(s), and whose Figures therefore display two-year recidivism rates that are biased upward. While the *relative* patterns they show (e.g. across race or sex) remain qualitatively similar, the levels are off. First, in the next two Figures, I replicate Figure 2 in Corbett-Davies et al. (2017) and then Figure 1 in Chouldechova (2016). Both of these Figures show the two-year recidivism rate by COMPAS score decile by race (for African-Americans and Caucasians only), but do so in a different format and style.

Here, and in the remaining Figures, I show the original Figures using ProPublica's two-year dataset on the left-hand side panels, and then on the right-hand side panels what the Figures look like using the corrected

two-year dataset that drops everyone with a COMPAS screen date post April 1, 2014. I try to replicate Figures as closely as possible to what they look like in the original publications.[46]
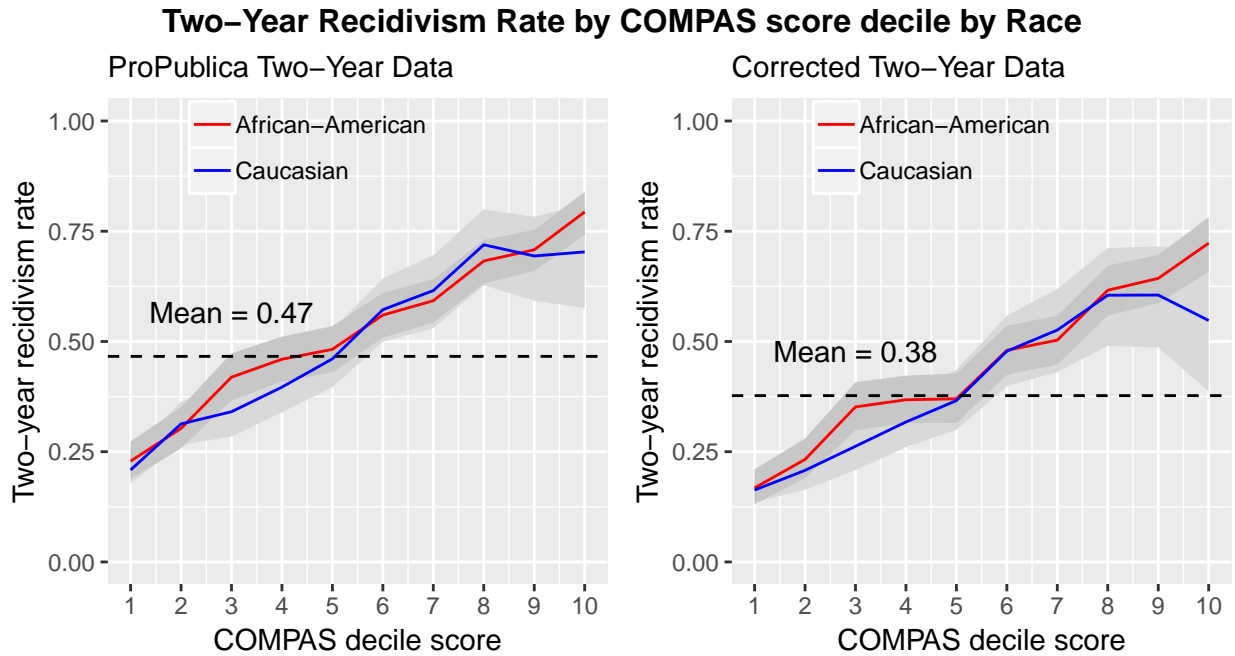
Two–Year Recidivism Rate by COMPAS score decile by Race

Figure 10: Two-Year Recidivism Rate by COMPAS score by Race (replicating Corbett-Davies et al. 2017)

**Two–Year Recidivism Rate by COMPAS score decile by Race**

Figure 11: Two-Year Recidivism Rate by COMPAS score by Race (replicating Chouldechova 2016)

---

[46]Using the same color schemes for example. Except for the following: Figure size or dimensions; the axis labels because I use a constant naming convention for axes in my paper for clarity; I add a *dashed line* for the *mean* recidivism rate; and finally, the sample sub-setting may not be exactly the same in every case, but it should be similar.

Next, I replicate a Figure by some of the same authors as the first Figure I replicated above, this time depicting the two-year recidivism rate by COMPAS score decile *by sex* (see Corbett-Davies and Goel (2018); this is Figure 1 in their paper). Since I have not previously shown the data breakdown by sex in this paper, I first include a Table with this information.

Table 26: Sex - ProPublica and Corrected Two-Year Data

|	ProPublica	Corrected
Female	1395	1213
Male	5819	5003
Total	7214	6216

Figure 12: Two-Year Recidivism Rate by COMPAS score by Sex (replicating Corbett-Davies and Goel 2018)

Finally, I replicate a Figure in the analysis done by DistrictDataLabs.[47] This Figure depicts the COMPAS score decile distribution by two-year recidivism status. (Figures are not numbered in their analysis)

---

[47]The full analysis by DistrictDataLabs can be found at https://www.districtdatalabs.com/fairness-and-bias-in-algorithms; last accessed on July 6, 2019.

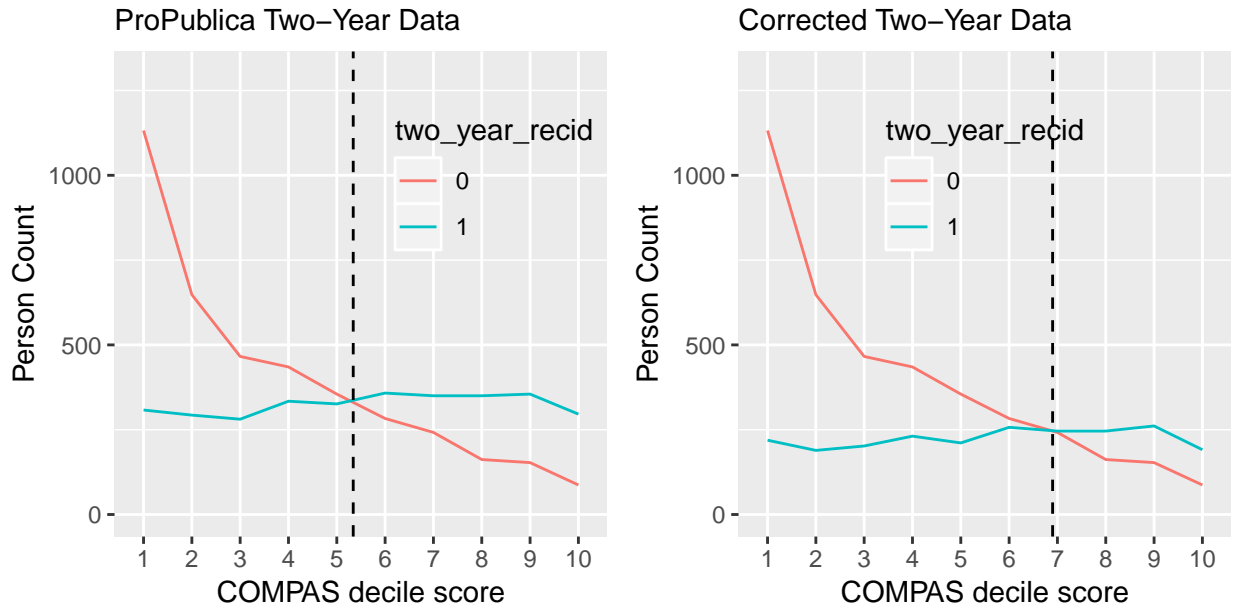**COMPAS Score Decile Distribution by Two–Year Recidivism Status**

Figure 13: COMPAS Score Decile Distribution by Two-Year Recidivism Status (replicating DistrictDataLabs)

In the last Figure I added a vertical dashed line in my paper (and in the previous Figures I added a horizontal dashed line). This vertical line in the last Figure is where the two curves in that Figure cross. That is the score at which there begin to be more recidivists than non-recidivists. This occurs at a (decile) average score slightly above 5 (around 5.34) in the ProPublica two-year dataset. But it occurs at a substantially higher (decile) average score of almost 7 (around 6.9) in the corrected two-year dataset. This is because there are fewer recidivists in the corrected data.[48]

---

[48]I plot these dashed vertical lines using a visual approximation since it is not clear how to obtain the exact crossing of the two curves given the discrete nature of the decile score data. But given that the difference between the two graphs is large, i.e. almost two score decile points, a visual approximation seems sufficient.